

notebook_Minh

April 30, 2023

```
{r setup, include=FALSE} title: "Untitled" author: "John Doe, Jane Doe"
```

```
[2]: """
Installation de Tensorflow et Keras est nécessaire pour lancer ce notebook
"""

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sns
from sklearn.linear_model import LinearRegression, Ridge, RidgeCV
from sklearn.kernel_ridge import KernelRidge
from sklearn.kernel_approximation import Nystroem
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense
from sklearn.metrics import r2_score
from sklearn.model_selection import cross_validate
%matplotlib inline
```

1 Exploration de données

```
[3]: df = pd.read_csv('kc_house_data.csv', index_col='id')
df
```

```
[3]:
```

	date	price	bedrooms	bathrooms	sqft_living	\
id						
7129300520	20141013T000000	221900.0	3	1.00	1180	
6414100192	20141209T000000	538000.0	3	2.25	2570	
5631500400	20150225T000000	180000.0	2	1.00	770	
2487200875	20141209T000000	604000.0	4	3.00	1960	
1954400510	20150218T000000	510000.0	3	2.00	1680	
...	
263000018	20140521T000000	360000.0	3	2.50	1530	
6600060120	20150223T000000	400000.0	4	2.50	2310	
1523300141	20140623T000000	402101.0	2	0.75	1020	
291310100	20150116T000000	400000.0	3	2.50	1600	

1523300157	20141015T000000	325000.0	2	0.75	1020		
------------	-----------------	----------	---	------	------	--	--


	sqft_lot	floors	waterfront	view	condition	grade	sqft_above \
id							
7129300520	5650	1.0	0	0	3	7	1180
6414100192	7242	2.0	0	0	3	7	2170
5631500400	10000	1.0	0	0	3	6	770
2487200875	5000	1.0	0	0	5	7	1050
1954400510	8080	1.0	0	0	3	8	1680
...
263000018	1131	3.0	0	0	3	8	1530
6600060120	5813	2.0	0	0	3	8	2310
1523300141	1350	2.0	0	0	3	7	1020
291310100	2388	2.0	0	0	3	8	1600
1523300157	1076	2.0	0	0	3	7	1020

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long \
id						
7129300520	0	1955	0	98178	47.5112	-122.257
6414100192	400	1951	1991	98125	47.7210	-122.319
5631500400	0	1933	0	98028	47.7379	-122.233
2487200875	910	1965	0	98136	47.5208	-122.393
1954400510	0	1987	0	98074	47.6168	-122.045
...
263000018	0	2009	0	98103	47.6993	-122.346
6600060120	0	2014	0	98146	47.5107	-122.362
1523300141	0	2009	0	98144	47.5944	-122.299
291310100	0	2004	0	98027	47.5345	-122.069
1523300157	0	2008	0	98144	47.5941	-122.299

	sqft_living15	sqft_lot15
id		
7129300520	1340	5650
6414100192	1690	7639
5631500400	2720	8062
2487200875	1360	5000
1954400510	1800	7503
...
263000018	1530	1509
6600060120	1830	7200
1523300141	1020	2007
291310100	1410	1287
1523300157	1020	1357

[21613 rows x 20 columns]

```
[4]: # Matrice de coorelation de DataFrame
cor = df.corr()
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(corr, annot=True, cmap=plt.cm.Reds, fmt='.2f',annot_kws={"fontsize":
↵8})
plt.show()
```



```
[5]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21613 entries, 7129300520 to 1523300157
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	date	21613 non-null	object
1	price	21613 non-null	float64
2	bedrooms	21613 non-null	int64
3	bathrooms	21613 non-null	float64
4	sqft_living	21613 non-null	int64
5	sqft_lot	21613 non-null	int64
6	floors	21613 non-null	float64
7	waterfront	21613 non-null	int64
8	view	21613 non-null	int64
9	condition	21613 non-null	int64
10	grade	21613 non-null	int64
11	sqft_above	21613 non-null	int64
12	sqft_basement	21613 non-null	int64
13	yr_built	21613 non-null	int64
14	yr_renovated	21613 non-null	int64
15	zipcode	21613 non-null	int64
16	lat	21613 non-null	float64
17	long	21613 non-null	float64
18	sqft_living15	21613 non-null	int64
19	sqft_lot15	21613 non-null	int64

dtypes: float64(5), int64(14), object(1)
memory usage: 3.5+ MB

2 Nettoyage de données

Nous avons supprimé les colonnes qui ne nous intéressent pas pour notre analyse.

```
[6]: new_df = df.drop(['date', 'zipcode', 'yr_renovated'], axis=1)
```

Nous avons décidé d'utiliser le `StandardScaler` de `sklearn` pour normaliser les données. Nous obtenons ainsi des données centrées réduites.

```
[7]: # Rescale data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
new_df['sqft_living'] = scaler.fit_transform(new_df['sqft_living'].values.
    ↪reshape(-1,1))
new_df['sqft_lot'] = scaler.fit_transform(new_df['sqft_lot'].values.
    ↪reshape(-1,1))
new_df['sqft_above'] = scaler.fit_transform(new_df['sqft_above'].values.
    ↪reshape(-1,1))
new_df['sqft_living15'] = scaler.fit_transform(new_df['sqft_living15'].values.
    ↪reshape(-1,1))
new_df['sqft_lot15'] = scaler.fit_transform(new_df['sqft_lot15'].values.
    ↪reshape(-1,1))
```

```

new_df['sqft_basement'] = scaler.fit_transform(new_df['sqft_basement'].values.
↳reshape(-1,1))
new_df['yr_built'] = scaler.fit_transform(new_df['yr_built'].values.
↳reshape(-1,1))
new_df['lat'] = scaler.fit_transform(new_df['lat'].values.reshape(-1,1))
new_df['long'] = scaler.fit_transform(new_df['long'].values.reshape(-1,1))
new_df.describe()

```

```

[7]:

```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
count	2.161300e+04	21613.000000	21613.000000	2.161300e+04	2.161300e+04	
mean	5.400881e+05	3.370842	2.114757	3.174253e-16	3.281921e-17	
std	3.671272e+05	0.930062	0.770163	1.000023e+00	1.000023e+00	
min	7.500000e+04	0.000000	0.000000	-1.948891e+00	-3.521759e-01	
25%	3.219500e+05	3.000000	1.750000	-7.108948e-01	-2.430487e-01	
50%	4.500000e+05	3.000000	2.250000	-1.849914e-01	-1.808075e-01	
75%	6.450000e+05	4.000000	2.500000	5.118578e-01	-1.066880e-01	
max	7.700000e+06	33.000000	8.000000	1.247807e+01	3.950434e+01	

	floors	waterfront	view	condition	grade	\
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000	
mean	1.494309	0.007542	0.234303	3.409430	7.656873	
std	0.539989	0.086517	0.766318	0.650743	1.175459	
min	1.000000	0.000000	0.000000	1.000000	1.000000	
25%	1.000000	0.000000	0.000000	3.000000	7.000000	
50%	1.500000	0.000000	0.000000	3.000000	7.000000	
75%	2.000000	0.000000	0.000000	4.000000	8.000000	
max	3.500000	1.000000	4.000000	5.000000	13.000000	

	sqft_above	sqft_basement	yr_built	lat	long	\
count	2.161300e+04	2.161300e+04	2.161300e+04	2.161300e+04	2.161300e+04	
mean	3.892022e-16	-2.022801e-15	3.592925e-15	-3.432958e-14	-3.663944e-14	
std	1.000023e+00	1.000023e+00	1.000023e+00	1.000023e+00	1.000023e+00	
min	-1.809494e+00	-6.586810e-01	-2.417383e+00	-2.916795e+00	-2.166543e+00	
25%	-7.226314e-01	-6.586810e-01	-6.810785e-01	-6.426977e-01	-8.102505e-01	
50%	-2.758102e-01	-6.586810e-01	1.360059e-01	8.478232e-02	-1.143518e-01	
75%	5.091458e-01	6.066704e-01	8.849999e-01	8.512345e-01	6.312541e-01	
max	9.204044e+00	1.023238e+01	1.497813e+00	1.570054e+00	6.383070e+00	

	sqft_living15	sqft_lot15
count	2.161300e+04	2.161300e+04
mean	-1.506632e-16	1.235382e-16
std	1.000023e+00	1.000023e+00
min	-2.316325e+00	-4.438052e-01
25%	-7.244971e-01	-2.808593e-01
50%	-2.138280e-01	-1.885636e-01
75%	5.448802e-01	-9.835556e-02
max	6.162239e+00	3.144029e+01

```
[8]: X = new_df.drop('price', axis=1)
     y = new_df['price']
```

2.1 Split des données en train et test

```
[9]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪ random_state=0)
     X_train
```

```
[9]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	\
id							
5100402668	3	1.00	-0.555193	-0.231701	1.0	0	
7856560480	3	2.50	-0.326539	-0.099155	1.0	0	
2872900010	3	1.50	-1.077829	-0.126630	1.0	0	
3216900070	4	2.50	0.141657	-0.193821	2.0	0	
976000790	3	2.50	-0.304762	-0.249736	2.0	0	
...	
2322069010	5	5.00	2.047104	1.906878	2.0	0	
2114700368	2	2.50	-0.740293	-0.334262	2.0	0	
5469501200	3	2.25	0.304981	-0.003790	1.0	0	
3751602797	4	2.00	0.315869	1.486207	2.0	0	
4038600260	4	2.25	0.326757	0.027258	1.0	0	

	view	condition	grade	sqft_above	sqft_basement	yr_built	\
id							
5100402668	0	4	7	-0.867546	0.471097	-1.055576	
7856560480	0	4	8	-0.698479	0.629266	0.306232	
2872900010	0	3	8	-0.843394	-0.658681	0.544548	
3216900070	0	3	8	0.509146	-0.658681	0.748819	
976000790	0	3	7	-0.662250	0.606670	0.476458	
...	
2322069010	0	3	10	2.622489	-0.658681	0.919045	
2114700368	0	3	8	-0.758860	-0.116388	1.259497	
5469501200	0	4	9	0.690290	-0.658681	0.238141	
3751602797	0	4	8	0.702366	-0.658681	0.238141	
4038600260	0	3	7	-0.299963	1.239346	-0.340627	

	lat	long	sqft_living15	sqft_lot15
id				
5100402668	0.968151	-0.746341	-0.315962	-0.233979
7856560480	-0.019143	0.460830	0.471927	-0.112383
2872900010	0.473060	1.263244	-0.403505	-0.106450
3216900070	-1.006438	0.219396	-0.024151	-0.211271
976000790	0.620288	-1.051685	-0.286781	-0.293202
...
2322069010	-1.299451	1.440769	0.369794	1.893510
2114700368	-0.186579	-0.959372	-1.351890	-0.411831

5469501200	-1.259035	0.396921	1.070140	0.059316
3751602797	-1.998784	-0.462301	0.180117	0.240465
4038600260	0.379239	0.666759	0.355203	-0.140768

[17290 rows x 16 columns]

3 Implementation de K-fold Cross Validation

```
[10]: def cross_validation(model, _x, _y, _cv=5):
    """
    Parameters
    -----
    model      : model to be used for cross validation
    _x          : features to train
    _y          : target to train
    _cv         : int
                  number of folds

    Returns
    -----
    results     : dictionary of average cross validation results

    """
    _scoring = ['r2', 'neg_mean_squared_error']
    results = cross_validate(model, _x, _y, cv=_cv, scoring=_scoring)
    results['test_mean_squared_error'] = -results['test_neg_mean_squared_error']
    results['mean_r2'] = results['test_r2'].mean()
    return results
```

4 Création des modèles

4.1 Linear regression model

```
[11]: linear_model = LinearRegression()
linear_model.fit(X_train, y_train)
```

```
[11]: LinearRegression()
```

```
[12]: res_validation = cross_validation(linear_model, X_train, y_train)
res_validation
```

```
[12]: {'fit_time': array([0.01049662, 0.06285238, 0.00599885, 0.00799727,
0.00498366]),
'score_time': array([0.00219822, 0.01600718, 0.00200152, 0.00101733,
0.00100088]),
'test_r2': array([0.67505938, 0.66494874, 0.6991278 , 0.71386052, 0.70856267]),
```

```

'test_neg_mean_squared_error': array([-5.80908872e+10, -3.82063063e+10,
-4.63087741e+10, -3.57464713e+10,
-3.54090643e+10]),
'test_mean_squared_error': array([5.80908872e+10, 3.82063063e+10,
4.63087741e+10, 3.57464713e+10,
3.54090643e+10]),
'mean_r2': 0.6923118242624663}

```

4.2 Prediction of Linear Regression Model

```

[13]: # Data scatter of predicted values
y_pred = linear_model.predict(X_test)
y_test = np.array(y_test)
y_pred, y_test

```

```

[13]: (array([ 379824.39845475, 1520132.09686934,  535192.63307498, ...,
348266.82410121, 202975.57538713, 403521.86733752]),
array([ 297000., 1578000.,  562100., ..., 369950., 300000., 575950.]))

```

```

[14]: # p-value on test set
r2_score(y_test, y_pred)

```

```

[14]: 0.6884620663104776

```

4.3 Ridge regression

4.3.1 Intuition de Ridge Regression

{python} Il faut la régularisation pour la Regression Ridge, même s'il faut que alpha ne soit pas tout petit. Sinon, il risque que les colonnes sont très corrélées. Si alpha est trop grand, la p-valeur ne descend pas beaucoup. En conclusion, nous pouvons dire que ce jeu de données est adapté pour la Regression Ridge.

```

[15]: ridge_reg = Ridge(alpha=1e8, solver='svd')

```

```

[16]: ridge_reg.fit(X_train, y_train)

```

```

[16]: Ridge(alpha=100000000.0, solver='svd')

```

```

[17]: # p-value de Ridge Regression
ridge_res_validation = cross_validation(ridge_reg, X_train, y_train)
ridge_res_validation['mean_r2']

```

```

[17]: -9.473387491352181e-05

```


4.3.2 Perform Leave-One-Out Cross Validation

```
[18]: ridgeCV = RidgeCV(alphas=[1e-6, 1e-3, 1, 1000], cv=5)

[19]: ridgeCV.fit(X_train, y_train)

[19]: RidgeCV(alphas=array([1.e-06, 1.e-03, 1.e+00, 1.e+03]), cv=5)

Meilleur  $\alpha$ : 1.0

[20]: ridgeCV.alpha_

[20]: 1.0

[21]: ridgeCV.score(X_test, y_test)

[21]: 0.6885181988335431
```

4.4 Prediction of Ridge Regression Model

```
[21]: def ridge_score(alpha=1.0):
        ridge_reg = Ridge(alpha=alpha, solver="svd")
        ridge_reg.fit(X_train, y_train)
        r2_ridge = cross_validation(ridge_reg, X_train, y_train)['mean_r2']
        return r2_ridge

[22]: ## Iterate over alpha values
alpha_values = np.linspace(0.01, 1, 100)
r2_ridge = []
for alpha in alpha_values:
    r2_ridge.append(ridge_score(alpha))
```

Nous pouvons voir que la R^2 ne varie pas beaucoup si nous faisons varier α . Cela veut dire que le modèle est assez stable.

```
[23]: r2_ridge

[23]: [0.6923119789388008,
        0.6923121331338864,
        0.69231228684803,
        0.6923124400815378,
        0.6923125928347165,
        0.6923127451078722,
        0.6923128969013108,
        0.6923130482153381,
        0.6923131990502597,
        0.6923133494063809,
        0.692313499284007,
        0.6923136486834429,
```

0.6923137976049933,
0.6923139460489631,
0.692314094015656,
0.692314241505377,
0.6923143885184295,
0.6923145350551175,
0.6923146811157446,
0.6923148267006141,
0.6923149718100292,
0.692315116444293,
0.6923152606037084,
0.6923154042885775,
0.6923155474992031,
0.6923156902358875,
0.6923158324989324,
0.6923159742886396,
0.692316115605311,
0.6923162564492479,
0.6923163968207513,
0.6923165367201225,
0.6923166761476625,
0.6923168151036714,
0.6923169535884497,
0.6923170916022983,
0.6923172291455166,
0.6923173662184047,
0.6923175028212623,
0.6923176389543885,
0.692317774618083,
0.6923179098126446,
0.6923180445383723,
0.6923181787955649,
0.6923183125845206,
0.6923184459055378,
0.6923185787589148,
0.6923187111449491,
0.6923188430639389,
0.6923189745161814,
0.692319105501974,
0.6923192360216139,
0.6923193660753978,
0.6923194956636228,
0.6923196247865852,
0.6923197534445814,
0.6923198816379076,
0.6923200093668598,
0.6923201366317338,

0.6923202634328252,
0.6923203897704291,
0.6923205156448411,
0.6923206410563562,
0.6923207660052688,
0.692320890491874,
0.692321014516466,
0.6923211380793391,
0.6923212611807876,
0.692321383821105,
0.692321506000585,
0.6923216277195212,
0.6923217489782071,
0.6923218697769354,
0.6923219901159994,
0.6923221099956915,
0.6923222294163044,
0.6923223483781304,
0.6923224668814619,
0.6923225849265905,
0.6923227025138081,
0.6923228196434063,
0.6923229363156764,
0.69232305253091,
0.6923231682893976,
0.6923232835914305,
0.6923233984372988,
0.6923235128272935,
0.6923236267617046,
0.6923237402408222,
0.6923238532649361,
0.6923239658343363,
0.6923240779493118,
0.6923241896101524,
0.6923243008171472,
0.692324411570585,
0.6923245218707546,
0.6923246317179446,
0.6923247411124432,
0.692324850054539,
0.6923249585445197]

4.5 Kernel Ridge Regression

```
[22]: feature_map_nystroem = Nystroem(gamma=.2,
                                     random_state=1,
                                     n_components=16)
feature_map_nystroem
```

```
[22]: Nystroem(gamma=0.2, n_components=16, random_state=1)
```

```
[25]: X_train_transformed = feature_map_nystroem.fit_transform(X_train)
X_train_transformed.shape
```

```
[25]: (17290, 16)
```

```
[28]: kernel_ridge = KernelRidge(alpha=1, kernel='rbf', gamma=.2)
```

```
[29]: kernel_ridge.fit(X_train_transformed, y_train)
```

```
[29]: KernelRidge(gamma=0.2, kernel='rbf')
```

```
[30]: kernel_ridge.score(X_train_transformed, y_train)
```

```
[30]: 0.37826677245632423
```

Nous avons décidé d'implémenter le Deep Neural Network pour obtenir un meilleur modèle.

4.6 Neural Network

```
[24]: neural_network = Sequential(
    [
        Dense(units=128, kernel_initializer='normal', input_dim=X_train.
↪shape[1], activation="relu"),
        Dense(units=256, kernel_initializer='normal', activation="relu"),
        Dense(units=256, kernel_initializer='normal', activation="relu"),
        Dense(units=1, activation="linear"),
    ], name="kc_model"
)
```

```
[25]: neural_network.summary()
```

Model: "kc_model"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	2176
dense_1 (Dense)	(None, 256)	33024
dense_2 (Dense)	(None, 256)	65792

dense_3 (Dense) (None, 1) 257

```
=====
Total params: 101,249
Trainable params: 101,249
Non-trainable params: 0
-----
```

```
[26]: #### Examine Weights shapes
[layer1, layer2, layer3, layer4] = neural_network.layers
W1,b1 = layer1.get_weights()
W2,b2 = layer2.get_weights()
W3,b3 = layer3.get_weights()
W4,b4 = layer4.get_weights()
print(f"W1 shape = {W1.shape}, b1 shape = {b1.shape}")
print(f"W2 shape = {W2.shape}, b2 shape = {b2.shape}")
print(f"W3 shape = {W3.shape}, b3 shape = {b3.shape}")
print(f"W4 shape = {W4.shape}, b4 shape = {b4.shape}")
```

```
W1 shape = (16, 128), b1 shape = (128,)
W2 shape = (128, 256), b2 shape = (256,)
W3 shape = (256, 256), b3 shape = (256,)
W4 shape = (256, 1), b4 shape = (1,)
```

```
[28]: X_train_numpy = np.array(X_train)
y_train_numpy = np.array(y_train)
X_train_numpy.shape
```

```
[28]: (17290, 16)
```

```
[29]: # Compile and fit the model
neural_network.compile(
    loss=tf.keras.losses.MeanAbsoluteError(),
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    metrics=[tf.keras.metrics.MeanAbsoluteError()]
)

history = neural_network.fit(
    X_train_numpy, y_train_numpy,
    epochs=50,
    batch_size=50,
    validation_split=0.2
)
```

```
Epoch 1/50
277/277 [=====] - 1s 2ms/step - loss: 361758.7500 -
mean_absolute_error: 361758.7500 - val_loss: 155124.0938 -
val_mean_absolute_error: 155124.0938
```

Epoch 2/50
277/277 [=====] - 0s 2ms/step - loss: 141654.6094 -
mean_absolute_error: 141654.6094 - val_loss: 124997.5156 -
val_mean_absolute_error: 124997.5156
Epoch 3/50
277/277 [=====] - 0s 2ms/step - loss: 126970.5391 -
mean_absolute_error: 126970.5391 - val_loss: 120336.3203 -
val_mean_absolute_error: 120336.3203
Epoch 4/50
277/277 [=====] - 0s 2ms/step - loss: 123883.9844 -
mean_absolute_error: 123883.9844 - val_loss: 117404.7734 -
val_mean_absolute_error: 117404.7734
Epoch 5/50
277/277 [=====] - 0s 2ms/step - loss: 121623.1641 -
mean_absolute_error: 121623.1641 - val_loss: 115586.2656 -
val_mean_absolute_error: 115586.2656
Epoch 6/50
277/277 [=====] - 0s 2ms/step - loss: 120109.9141 -
mean_absolute_error: 120109.9141 - val_loss: 114070.8047 -
val_mean_absolute_error: 114070.8047
Epoch 7/50
277/277 [=====] - 0s 2ms/step - loss: 118616.9297 -
mean_absolute_error: 118616.9297 - val_loss: 112459.3359 -
val_mean_absolute_error: 112459.3359
Epoch 8/50
277/277 [=====] - 0s 2ms/step - loss: 117205.8750 -
mean_absolute_error: 117205.8750 - val_loss: 112344.6250 -
val_mean_absolute_error: 112344.6250
Epoch 9/50
277/277 [=====] - 0s 2ms/step - loss: 116026.7969 -
mean_absolute_error: 116026.7969 - val_loss: 110621.2578 -
val_mean_absolute_error: 110621.2578
Epoch 10/50
277/277 [=====] - 0s 2ms/step - loss: 114882.4219 -
mean_absolute_error: 114882.4219 - val_loss: 109255.5938 -
val_mean_absolute_error: 109255.5938
Epoch 11/50
277/277 [=====] - 0s 2ms/step - loss: 113941.2109 -
mean_absolute_error: 113941.2109 - val_loss: 110929.6172 -
val_mean_absolute_error: 110929.6172
Epoch 12/50
277/277 [=====] - 0s 2ms/step - loss: 113460.0547 -
mean_absolute_error: 113460.0547 - val_loss: 107895.9688 -
val_mean_absolute_error: 107895.9688
Epoch 13/50
277/277 [=====] - 0s 2ms/step - loss: 112663.3047 -
mean_absolute_error: 112663.3047 - val_loss: 107421.8125 -
val_mean_absolute_error: 107421.8125

Epoch 14/50
277/277 [=====] - 0s 2ms/step - loss: 112046.0859 -
mean_absolute_error: 112046.0859 - val_loss: 107013.5000 -
val_mean_absolute_error: 107013.5000

Epoch 15/50
277/277 [=====] - 0s 2ms/step - loss: 111549.4531 -
mean_absolute_error: 111549.4531 - val_loss: 107145.0156 -
val_mean_absolute_error: 107145.0156

Epoch 16/50
277/277 [=====] - 0s 2ms/step - loss: 111215.2266 -
mean_absolute_error: 111215.2266 - val_loss: 107112.7891 -
val_mean_absolute_error: 107112.7891

Epoch 17/50
277/277 [=====] - 0s 2ms/step - loss: 110785.0938 -
mean_absolute_error: 110785.0938 - val_loss: 106216.1562 -
val_mean_absolute_error: 106216.1562

Epoch 18/50
277/277 [=====] - 0s 2ms/step - loss: 110588.5000 -
mean_absolute_error: 110588.5000 - val_loss: 106097.2109 -
val_mean_absolute_error: 106097.2109

Epoch 19/50
277/277 [=====] - 0s 2ms/step - loss: 110157.4609 -
mean_absolute_error: 110157.4844 - val_loss: 105967.0469 -
val_mean_absolute_error: 105967.0469

Epoch 20/50
277/277 [=====] - 0s 2ms/step - loss: 109995.2031 -
mean_absolute_error: 109995.1875 - val_loss: 106093.3125 -
val_mean_absolute_error: 106093.3125

Epoch 21/50
277/277 [=====] - 0s 2ms/step - loss: 109714.9219 -
mean_absolute_error: 109714.9219 - val_loss: 105784.0781 -
val_mean_absolute_error: 105784.0781

Epoch 22/50
277/277 [=====] - 0s 2ms/step - loss: 109561.1406 -
mean_absolute_error: 109561.1406 - val_loss: 105901.7266 -
val_mean_absolute_error: 105901.7266

Epoch 23/50
277/277 [=====] - 0s 2ms/step - loss: 109206.2266 -
mean_absolute_error: 109206.2266 - val_loss: 105472.8359 -
val_mean_absolute_error: 105472.8359

Epoch 24/50
277/277 [=====] - 0s 2ms/step - loss: 108979.5156 -
mean_absolute_error: 108979.5156 - val_loss: 105321.5703 -
val_mean_absolute_error: 105321.5703


Epoch 25/50
277/277 [=====] - 0s 2ms/step - loss: 108751.5781 -
mean_absolute_error: 108751.5781 - val_loss: 105163.6484 -
val_mean_absolute_error: 105163.6484

Epoch 26/50
277/277 [=====] - 0s 2ms/step - loss: 108611.1406 -
mean_absolute_error: 108611.1406 - val_loss: 105323.1953 -
val_mean_absolute_error: 105323.1953
Epoch 27/50
277/277 [=====] - 0s 2ms/step - loss: 108587.8906 -
mean_absolute_error: 108587.8906 - val_loss: 104946.6172 -
val_mean_absolute_error: 104946.6172
Epoch 28/50
277/277 [=====] - 0s 2ms/step - loss: 108280.2109 -
mean_absolute_error: 108280.2109 - val_loss: 105723.8750 -
val_mean_absolute_error: 105723.8750
Epoch 29/50
277/277 [=====] - 1s 2ms/step - loss: 108220.4688 -
mean_absolute_error: 108220.4688 - val_loss: 104851.5234 -
val_mean_absolute_error: 104851.5234
Epoch 30/50
277/277 [=====] - 0s 2ms/step - loss: 108115.5625 -
mean_absolute_error: 108115.5625 - val_loss: 104801.7500 -
val_mean_absolute_error: 104801.7500
Epoch 31/50
277/277 [=====] - 1s 2ms/step - loss: 107901.3047 -
mean_absolute_error: 107901.3047 - val_loss: 104851.4219 -
val_mean_absolute_error: 104851.4219
Epoch 32/50
277/277 [=====] - 1s 2ms/step - loss: 107854.3438 -
mean_absolute_error: 107854.3438 - val_loss: 104605.5625 -
val_mean_absolute_error: 104605.5625
Epoch 33/50
277/277 [=====] - 0s 2ms/step - loss: 107688.7266 -
mean_absolute_error: 107688.7266 - val_loss: 105052.9297 -
val_mean_absolute_error: 105052.9297
Epoch 34/50
277/277 [=====] - 0s 2ms/step - loss: 107713.3906 -
mean_absolute_error: 107713.3906 - val_loss: 104819.6484 -
val_mean_absolute_error: 104819.6484
Epoch 35/50
277/277 [=====] - 0s 2ms/step - loss: 107379.9453 -
mean_absolute_error: 107379.9453 - val_loss: 104927.8359 -
val_mean_absolute_error: 104927.8359
Epoch 36/50
277/277 [=====] - 0s 2ms/step - loss: 107327.0156 -
mean_absolute_error: 107327.0156 - val_loss: 104815.0391 -
val_mean_absolute_error: 104815.0391
Epoch 37/50
277/277 [=====] - 0s 2ms/step - loss: 107233.3906 -
mean_absolute_error: 107233.3906 - val_loss: 104432.8594 -
val_mean_absolute_error: 104432.8594

Epoch 38/50
277/277 [=====] - 0s 2ms/step - loss: 107128.2656 -
mean_absolute_error: 107128.2656 - val_loss: 104360.8672 -
val_mean_absolute_error: 104360.8672
Epoch 39/50
277/277 [=====] - 0s 2ms/step - loss: 106943.2500 -
mean_absolute_error: 106943.2500 - val_loss: 104153.7422 -
val_mean_absolute_error: 104153.7422
Epoch 40/50
277/277 [=====] - 0s 2ms/step - loss: 106985.8750 -
mean_absolute_error: 106985.8750 - val_loss: 104217.3281 -
val_mean_absolute_error: 104217.3281
Epoch 41/50
277/277 [=====] - 0s 2ms/step - loss: 106825.0859 -
mean_absolute_error: 106825.0859 - val_loss: 104136.8672 -
val_mean_absolute_error: 104136.8672
Epoch 42/50
277/277 [=====] - 0s 2ms/step - loss: 106819.1484 -
mean_absolute_error: 106819.1484 - val_loss: 103906.4375 -
val_mean_absolute_error: 103906.4375
Epoch 43/50
277/277 [=====] - 0s 2ms/step - loss: 106654.2969 -
mean_absolute_error: 106654.2969 - val_loss: 104109.7891 -
val_mean_absolute_error: 104109.7891
Epoch 44/50
277/277 [=====] - 0s 2ms/step - loss: 106452.1953 -
mean_absolute_error: 106452.1953 - val_loss: 104054.7500 -
val_mean_absolute_error: 104054.7500
Epoch 45/50
277/277 [=====] - 0s 2ms/step - loss: 106468.8359 -
mean_absolute_error: 106468.8359 - val_loss: 103671.9609 -
val_mean_absolute_error: 103671.9609
Epoch 46/50
277/277 [=====] - 0s 2ms/step - loss: 106318.0859 -
mean_absolute_error: 106318.0859 - val_loss: 103678.5391 -
val_mean_absolute_error: 103678.5391
Epoch 47/50
277/277 [=====] - 0s 2ms/step - loss: 106125.0234 -
mean_absolute_error: 106125.0234 - val_loss: 103490.8594 -
val_mean_absolute_error: 103490.8672
Epoch 48/50
277/277 [=====] - 0s 2ms/step - loss: 106151.3828 -
mean_absolute_error: 106151.3828 - val_loss: 103520.8672 -
val_mean_absolute_error: 103520.8672
Epoch 49/50
277/277 [=====] - 1s 2ms/step - loss: 106128.6484 -
mean_absolute_error: 106128.6484 - val_loss: 103634.1016 -
val_mean_absolute_error: 103634.1016

```
Epoch 50/50
277/277 [=====] - 0s 2ms/step - loss: 105923.1953 -
mean_absolute_error: 105923.1953 - val_loss: 103399.5391 -
val_mean_absolute_error: 103399.5391
```

```
[30]: # Visualisation
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Val Loss')
plt.xlabel('Epoch')
plt.legend(['Train'], loc='upper left')
plt.show()
```



4.7 Prediction of Neural Network Model

```
[31]: X_test_numpy = np.array(X_test)
y_test_numpy = np.array(y_test)
```

```
[32]: # Data scatter of predicted values
y_pred = neural_network.predict(X_test_numpy)
y_pred = np.reshape(y_pred, (y_pred.shape[0],))
```

```
y_test_numpy, y_pred
```

```
136/136 [=====] - 0s 805us/step
```

```
[32]: (array([ 297000., 1578000., 562100., ..., 369950., 300000., 575950.]),  
      array([ 438204.66, 1385815.5 , 560890. , ..., 401254.9 , 251818.45,  
            390145.7 ], dtype=float32))
```

```
[33]: # p-value  
      r2_score(y_test_numpy, y_pred)
```

```
[33]: 0.7394895122908761
```

4.8 Implementation des Neural Network modèles avec des différentes learning rates

```
[34]: def neural_network(num_layer, units, learning_rate=0.001):  
      """  
      Parameters  
      -----  
      num_layer      : int  
                      number of layers  
      units          : list  
                      number of units in each layer  
      learning_rate  : float  
                      learning rate of the optimizer of the model  
      """  
      if (len(units) != num_layer):  
          raise ValueError("Number of list of units must be equal to number of_  
↪layers")  
      model = Sequential()  
      model.add(Dense(units=units[0], kernel_initializer='normal',_  
↪input_dim=X_train.shape[1], activation="relu"))  
      for i in range(1, num_layer-1):  
          model.add(Dense(units=units[i], kernel_initializer='normal',_  
↪activation="relu"))  
      model.add(Dense(units=1, activation="linear"))  
      model.compile(  
          loss=tf.keras.losses.MeanAbsoluteError(),  
          optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),  
          metrics=[tf.keras.metrics.MeanAbsoluteError()]  
      )  
      return model
```

```
[35]: # Itération  
      histories_neural_network = []  
      def neural_network_score(num_layer, units, learning_rate=[0.001, 0.01, 0.1]):  
          """
```

```

Parameters
-----
num_layer      : int
                  number of layers
units          : list
                  number of units in each layer
learning_rate  : list
                  learning rate of the optimizer of the model
"""
if (len(units) != num_layer):
    raise ValueError("Number of list of units must be equal to number of layers")
for i in range(len(learning_rate)):
    model = neural_network(num_layer, units, learning_rate=learning_rate[i])
    model.save(f"neural_network_model_{i}.h5")
    history = model.fit(
        X_train_numpy, y_train_numpy,
        epochs=50,
        batch_size=50,
        validation_split=0.2
    )
    histories_neural_network.append(history)
    y_pred = model.predict(X_test)
    y_pred = np.reshape(y_pred, (y_pred.shape[0],))
    print(f"R2 score on test set with learning rate {learning_rate[i]} is {r2_score(y_test_numpy, y_pred)}")
    print(f"Mean absolute error on test set with learning rate {learning_rate[i]} is {history.history['val_loss'][-1]}")

```

```

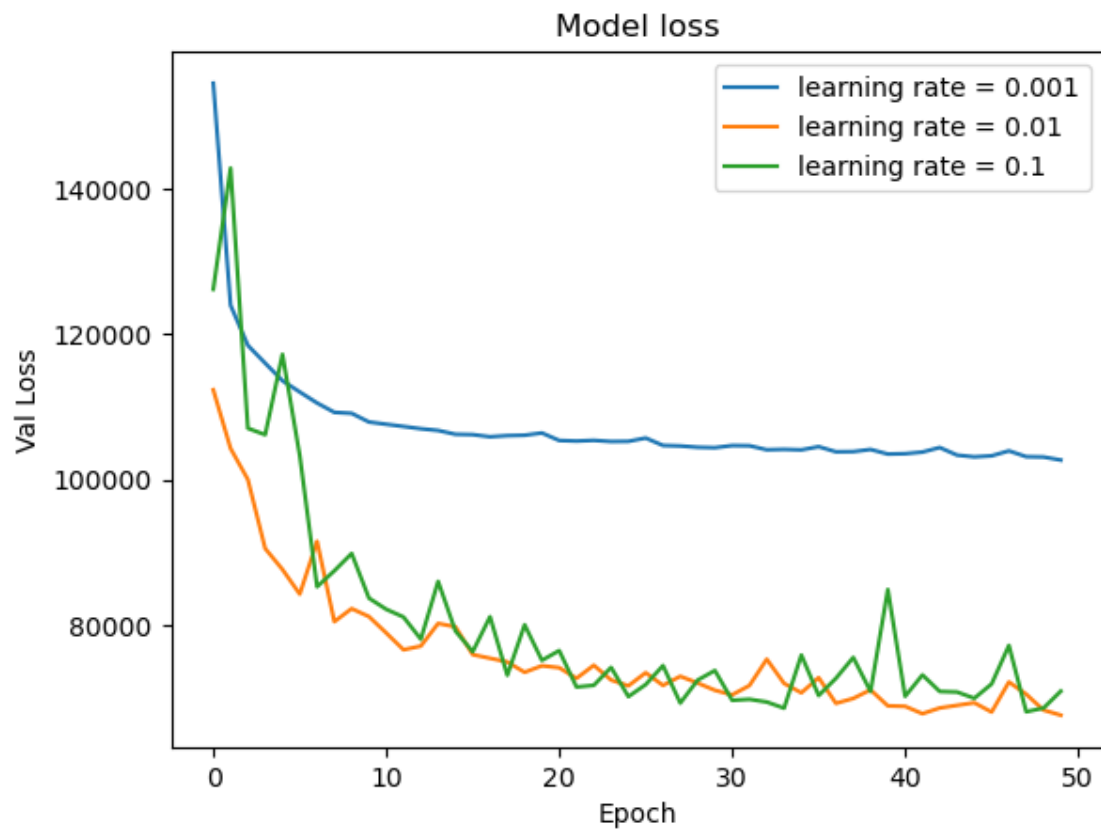
[ ]: num_layer = 4
     units = [128, 256, 256, 1]
     neural_network_score(num_layer, units)

```

```

[37]: # Visualisation
      plt.plot(histories_neural_network[0].history['val_loss'])
      plt.plot(histories_neural_network[1].history['val_loss'])
      plt.plot(histories_neural_network[2].history['val_loss'])
      plt.title('Model loss')
      plt.ylabel('Val Loss')
      plt.xlabel('Epoch')
      plt.legend(["learning rate = 0.001", "learning rate = 0.01", "learning rate = 0.1"], loc='upper right')
      plt.show()

```



Nous pouvons voir que avec `learning rate = 0.01`, le modèle converge plus vite que le modèle avec `learning rate = 0.001`. Il est plus stable que le modèle avec `learning rate = 0.1`. Nous choisissons donc le modèle avec `learning rate = 0.01` comme modèle plus optimisé.